


Building an AI Agent from Scratch – A Practical Guide

▼ Type

@datasciencebrain



AI agents are no longer just research experiments. They're powering customer support, automating workflows, driving apps, and even assisting in writing code. If you're looking to build an AI agent from the ground up—without relying on heavy pre-built solutions—this guide walks you through every essential concept and step.

Let's begin from scratch and work our way up to building scalable and functional AI agents for real-world applications.

1. What is an AI Agent?

An AI agent is a system that:

- Perceives its environment (through inputs),
- Makes decisions (via logic, rules, or models),
- Acts upon those decisions (produces structured output, calls APIs, controls other systems).

In simple terms, it's like giving AI a brain (for thinking) and hands (for acting).

2. Core Components of an AI Agent

To build an agent, you typically need to implement these layers:

1. **Input Handler** – Accepts natural language input or system signals.
 2. **Memory / State** – Stores conversation history or relevant past data.
 3. **Reasoning Engine** – The logic or model that decides the next step.
 4. **Tools or Actions** – Interfaces to perform tasks like fetching data, calling APIs, etc.
 5. **Output Generator** – Converts internal decisions into human-readable responses or system actions.
-

3. Step-by-Step Roadmap to Build an AI Agent

Step 1: Define the Agent's Goal

Start by clearly defining what you want the agent to do. For example:

- Book a meeting.
- Retrieve company data from a database.
- Assist in coding tasks.
- Help answer domain-specific queries.

Avoid vague goals like “chatbot to answer everything.” Focus makes everything easier to design, test, and scale.

Step 2: Choose Your Language Model

For beginners, using pre-trained language models via APIs is the best route.

Options:

- OpenAI (GPT-4)
- Anthropic (Claude)
- Google Gemini

You can start with `openai`'s API, as it's well-documented and supports function calling (aka tool use).

Install the SDK:

```
pip install openai
```

Basic setup:

```
import openai

openai.api_key = "your-api-key"

response = openai.ChatCompletion.create(
    model="gpt-4",
    messages=[{"role": "user", "content": "What's the weather in Kochi?"}]
)
print(response['choices'][0]['message']['content'])
```

Step 3: Add Structured Output (Tool Use)

Raw text is limiting. To do meaningful work, agents must return structured output or trigger tools.

Use OpenAI's function calling to describe tasks in structured JSON format:

```
functions = [
    {
```

```

    "name": "get_weather",
    "description": "Fetches the weather forecast",
    "parameters": {
      "type": "object",
      "properties": {
        "location": {"type": "string", "description": "City name"}
      },
      "required": ["location"]
    }
  }
]

```

You pass this into the API call, and it will return:

```

{
  "function_call": {
    "name": "get_weather",
    "arguments": "{ \"location\": \"Kochi\" }"
  }
}

```

You then connect this to your weather API and return results.

Step 4: Build a Tool Layer (Actions)

This is where your agent “does” things. Each function should be modular.

Example:

```

def get_weather(location):
    # Connect to weather API like OpenWeatherMap
    return f"The current weather in {location} is 30°C and sunny."

```

Map function names to Python functions:

```

tools = {
  "get_weather": get_weather
}

```

```
}
```

When the model returns a function call, you extract it and invoke the tool.

Step 5: Maintain Memory or Context

For short tasks, storing memory in a simple Python list is fine:

```
conversation = [  
    {"role": "user", "content": "What's the weather in Kochi?"},  
    {"role": "assistant", "content": "It's 30°C and sunny."}  
]
```

For advanced cases:

- Use **Redis** or **ChromaDB** for vector search memory.
 - Log user context and tool outputs to dynamically feed into prompts.
-

Step 6: Design the Reasoning Strategy

Two approaches:

- **Single-step agents:** Get input → respond.
- **Multi-step agents:** Think before acting, possibly replan.

A common pattern is **ReAct** (Reason + Act):

1. Model reflects on the problem.
 2. Chooses a tool.
 3. Observes tool output.
 4. Repeats until solution is found.
-

Step 7: Implement Looping Agents (Multi-Action Agents)

This is how agents autonomously reach goals:

```

while not done:
    response = call_openai(conversation, functions)
    if response.function_call:
        function_name = response.function_call.name
        arguments = json.loads(response.function_call.arguments)
        result = tools[function_name](**arguments)
        conversation.append({"role": "function", "name": function_name, "content": result})
    else:
        print(response.content)
        done = True

```

This way, the model keeps getting updated with the latest context and continues reasoning until it produces a final response.

Step 8: Add Safeguards and Error Handling

Things will break. Your agent should:

- Handle failed tool calls.
- Retry API errors.
- Detect invalid JSON and correct it (you can even send back the error and ask the model to fix it).

Add timeouts, logging, and try-except blocks.

9. Optional Add-ons

- **UI Layer:** Build a frontend using Streamlit, React, or Gradio.
- **Authentication:** Add token-based auth if deployed online.
- **Analytics:** Log inputs, tool usage, and user feedback to improve performance.
- **Cost Control:** Set max tokens and batch requests where possible.

10. Best Practices

- Keep prompts deterministic; define rules clearly.
 - Log everything during development to trace failures.
 - Don't let agents run tools indefinitely—set a max loop count.
 - Make tools modular and test them independently.
 - Always validate model output before executing actions.
-

11. Common Challenges and Fixes

Challenge	Fix
Invalid JSON from model	Use regex fixers or request correction from the model
Tool call has missing params	Validate inputs before calling
Loop never ends	Set a loop cap or detect goal fulfillment
Memory too long	Use summary-based memory or vector memory
User input misunderstood	Include user role description and task context in system prompt

12. Final Thoughts

AI agents aren't just chatbots—they're systems capable of acting with purpose. Start with simple agents that call one or two tools, and as you grow more comfortable, add memory, goal reasoning, and advanced chaining.

What matters most isn't how complex your model is, but how thoughtfully you design the workflow between model, tools, and users.
